

Computer Vision + ML using GPU/CUDA

CuCNN [[GitHub](#)]

*CSN- 291: Computer Architecture and Microprocessor
(CAM)*

CuCNN

(Cuda Convolutional Neural Network)

- Jitesh Jain

1. Problem Statement

Implement a simple **3-layer CNN**[5][6] (Convolutional Neural Network) for Image Classification on **MNIST**[1][2] using **CUDA**[3][4] **Programming** and observe the *effect of different kernel settings on the performance of the model*.

2. Novelty

This work presents a study about the *effect of using different Grid Size and Block Size on a CNN model's performance*, which is supported by the experiments.

It also presents forward a *simple 3-layer CNN architecture* that attains a *test accuracy in excess of 97%* along with a *training time of only 5 minutes* on a **Tesla T4** (available for free on *GoogleColab*[7]) system.

3. Evaluation Parameters

We use the **MNIST**[1][2] dataset for training and testing of our CNN model. **MNIST**[1][2] is an extensive database of handwritten digits

(0-9) containing **60000 images of size 28x28**. Due to its *wide usage* for *testing and training of classification models in the field of Machine Learning*, we consider it as a **benchmark** for evaluating our model's performance.

We evaluate our model based on three parameters:

- **Accuracy:** *The ratio of the number of correctly classified images to the total number of images.*
- **Train Error:** *The total error calculated during the training as the sum of the euclidean norm (implemented in CUDA using `cublasSnmr2` in the `cuBLAS`[8] library) of the vector containing the sum of the differences between the target and predicted probability for each digit of a given image divided by the number of images.*
- **Training Time:** *The time taken to train the CNN model on a Tesla T4 GPU for 50 epochs.*

4. Methodology


In this section, we first provide an overview of CNN in **Section 4.1**. We briefly explain our 3-layer CNN architecture in **Section 4.2**, followed by introducing the terms in CUDA programming in **Section 4.3**. Finally, in **Section 4.4**, we present the implementation details.

4.1 Convolutional Neural Network (CNN)

A CNN[5][6] is a Neural Network specifically designed for computer vision applications like image classification, semantic segmentation, etc.

The basic building blocks of a CNN are:


1. **Convolution layer:** It consists of *sliding window* operations where a **filter** (also called the **kernel**, the *middle one in the image below*) slides over the **input feature map** (left one in the image below) and *outputs one value corresponding to each of the overlaps*. It works using *parameter sharing* (as only the kernels have the learnable parameters), thus reducing the parameters by a considerable margin compared to a *linear deep neural network*[9].



Convolution Operation

[Source]

- 2. Pooling Layer:** It is used to reduce the spatial size of the feature map output by a convolution layer. There are two types of pooling: Average Pooling and Max Pooling (see the image below). We don't use a pooling layer in our architecture because the input image already has a small size (28 x 28).



Pooling Operation

[[Source](#)]

- 3. Activation Function:** It is generally used to induce non-linearity in the network to make the model more robust to variations in the data. We use *sigmoid* as our activation function.

$$S(x) = 1 / (1 + e^x)$$

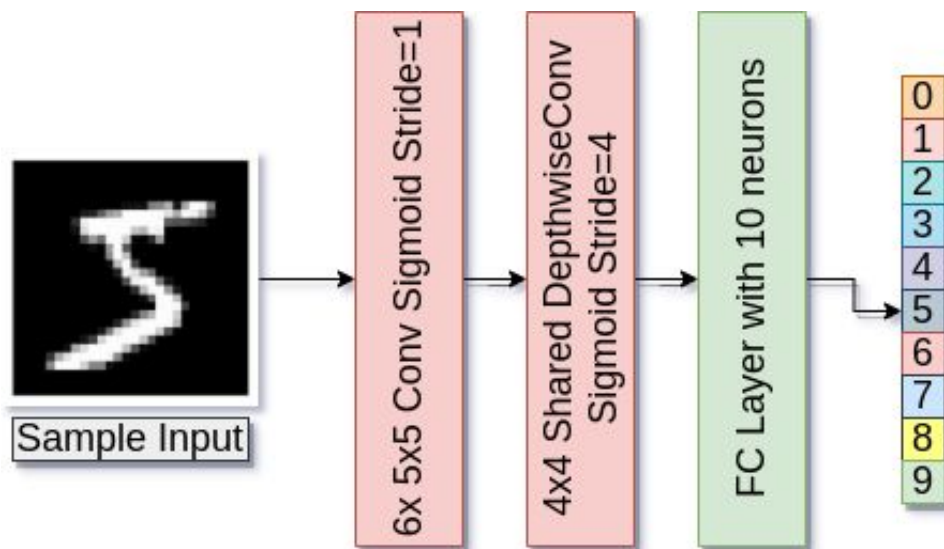
Sigmoid Function

- 4. Fully-Connected Layer:** It is a linear layer with multiple nodes at the end stages of the network to predict the probability of each category. We use an FC layer with 10 nodes in our architecture.

4.2 The CNN Architecture

The 3-Layer CNN consists of:

- **Convolution Layer:** Applies a **Convolution Operation** with 6 kernels of size 5×5 with $stride=1$ on the input image (size= 28×28) to output a map of shape= $24 \times 24 \times 6$.
- **Shared Depthwise Convolution Layer:** Applies a *shared Depthwise* (the same kernel applied to different channels of the input from the previous layer) *Convolution Operation* with a 4×4 kernel with $stride = 4$ on each channel of the previous Convolution layer's output feature map to output a map of shape = $6 \times 6 \times 6$.
- **Fully Connected Layer:** *Flattens* the output from the previous layer to a layer with 10 nodes, with each node's value representing the *probability* of a digit from 0-9.



3-layer CNN Architecture (made using draw.io)

4.3 CUDA Programming


CUDA[3][4] is a parallel computing platform and programming model developed by Nvidia for general computing on its GPUs. CUDA enables developers to speed up compute-intensive applications by harnessing GPUs' power for the computation's parallelizable part.

Some technical terms:

- **Host:** The *CPU* that makes calls to the kernels.
- **Device:** The *GPU*.
- **__global__ Function:** Called by the *host* and executed on the *device*.
- **__device__ function:** Called by the *device* and executed on the *device*.
- **Block Size:** Number of threads inside a block.
- **Grid Size:** Number of blocks inside a grid.

4.4 Implementation Details

We implemented specific feed-forward *and backpropagation* functions for our CNN model along with activations and error functions, all using the **CUDA** framework. Some important details are:



High-level Overview of CUDA

[Source]

- We used simple *float multi-dimensional arrays* to implement the *parameter structures* (weights and biases) and *outputs* for the CNN.

- The functions called during feedforward propagation and backpropagation were specified as **__global__** as they are called during training from the **host (CPU)**.
- We train our model for **50 epochs (or iterations)**, in different settings) and 100 epochs (or iterations, in one setting).
- The **best performing kernel sizes** are: **Grid Size = 64, Block Size = 64**, i.e, *kernel <<<64,64>>>*.

5. Results/Experiments

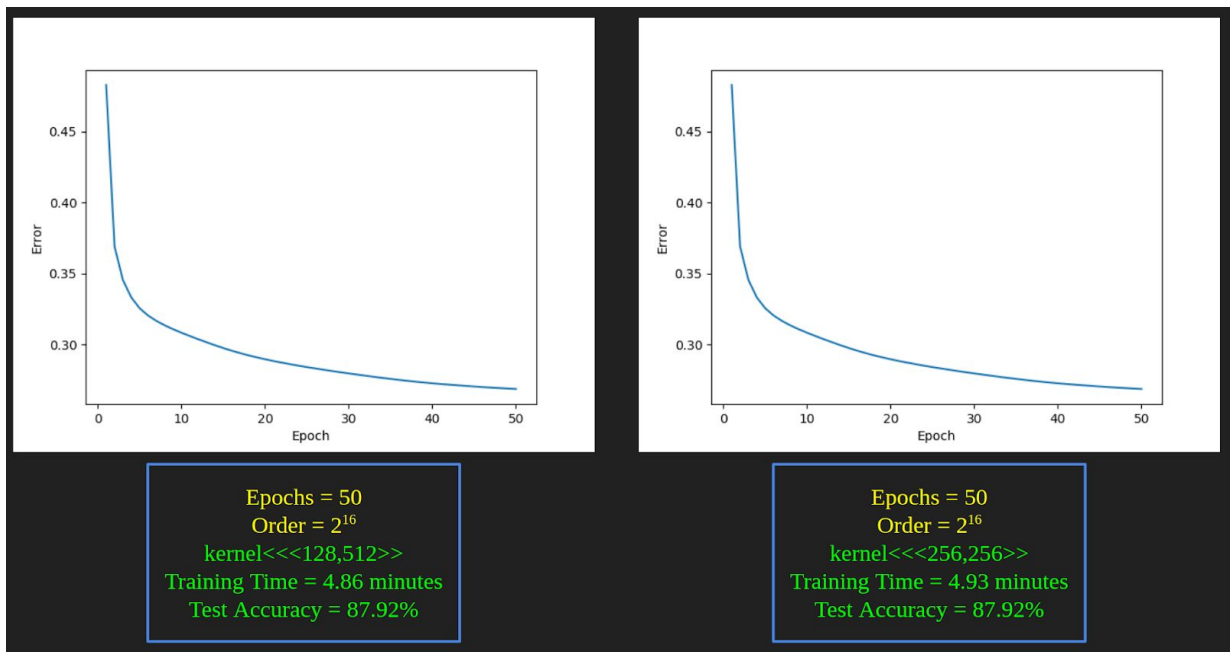
We perform experiments on *different Block Sizes* and *Grid Sizes* settings for the kernel and make *useful observations*. All experiments were performed on a **Tesla T4 GPU**.

Grid Size	Block Size	Order (2)	Epochs	Test Accuracy	Training Time (Minutes)
64	64	12	50	97.12	4.54
64	64	12	100	97.41	9.10
128	512	16	50	87.92	4.86
256	256	16	50	87.92	4.93
16	16	8	50	97.12	10.02

Results

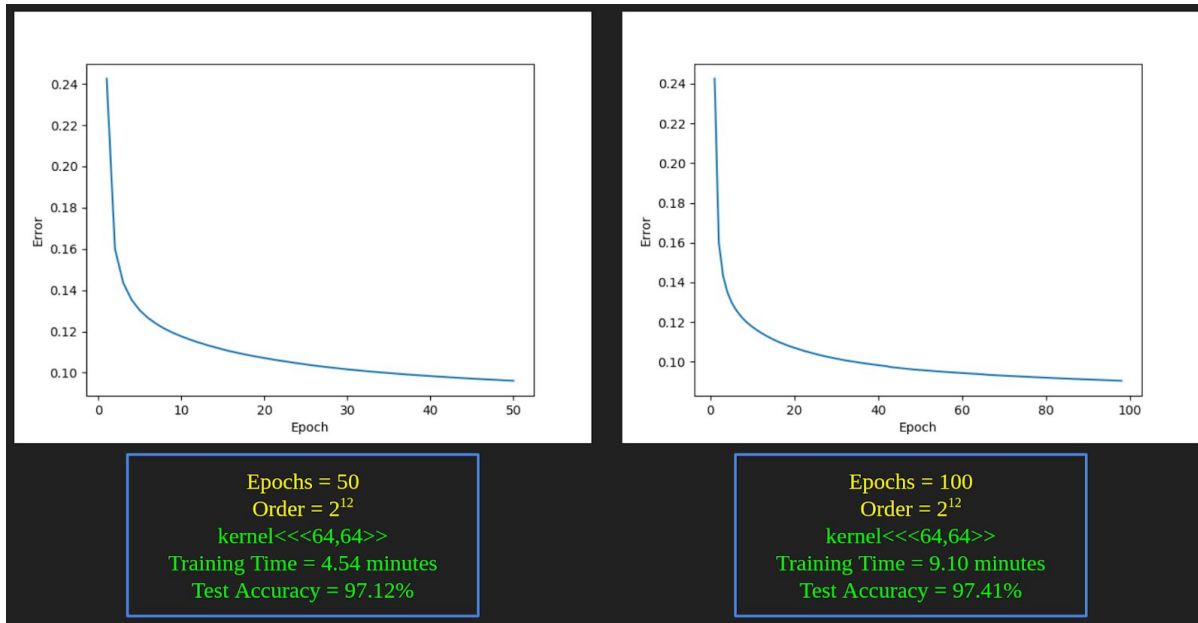
Observations

- The model's performance depends on the **Grid Size** (number of Blocks) and **Block Size** (number of Threads).
- Products of the same order give almost the same results (in terms of time and accuracy).



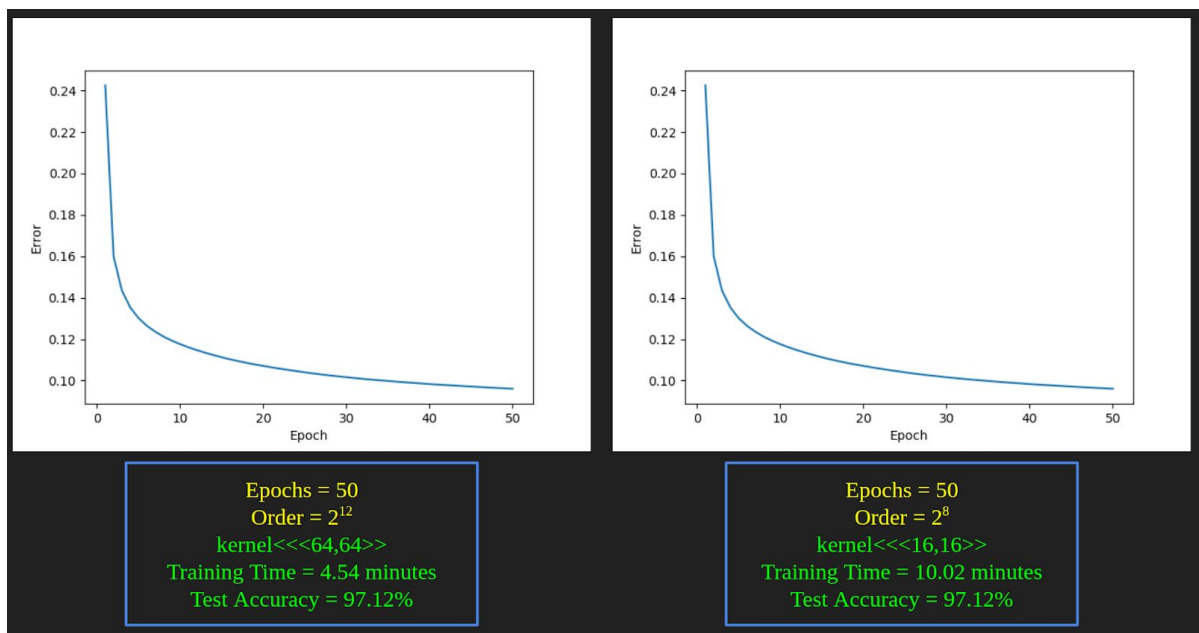
Similar Order (product) give similar results

- Higher-order (product, 2^{16}) gives *inaccurate results* compared to those of lower-order (2^{12}) of 2.



The training doesn't improve accuracy a lot after 50 epochs

- As the product's order *becomes smaller*, training time increases even if accuracy remains the same (2^8 order v/s 2^{12} order).



Training time increases although the accuracy is the same

6. Conclusion

Here, we present a simple 3-layer *CNN*[5][6] implemented using *CUDA*[3][4]. We call the resulting implementation ***CuCNN***. Along with successful training and inference on the *MNIST*[1][2] dataset, we also present some useful insights into the implementation of ML algorithms on CUDA.

Implementation and Report by:

Jitesh Jain (Enrolment Number: **19114039**)

References

1. [MNIST handwritten digit database](#)
2. <https://github.com/projectgalatea/mnist#mnist>
3. [Getting Started with CUDA](#)
4. [About CUDA](#)
5. [ImageNet Classification with Deep Convolutional Neural Networks](#)
6. <https://towardsdatascience.com/an-introduction-to-convolutional-neural-networks-eb0b60b58fd7?gi=31c7f978cb3b>
7. [Google Colab Introduction](#)
8. [cuBLAS Library](#)
9. [Neural networks and deep learning](#)
10. [CS334 on Udacity](#)